

HW3

March 18, 2018

1 HOMEWORK 3

1.1 PART 1

```
In [16]: import matplotlib.pyplot as plt
import numpy as np
import pandas as pd
import math

X_train = np.array(pd.read_csv("./gaussian_process/X_train.csv", header=None))
X_test = np.array(pd.read_csv("./gaussian_process/X_test.csv", header=None))
y_train = np.array(pd.read_csv("./gaussian_process/y_train.csv", header=None))
y_test = np.array(pd.read_csv("./gaussian_process/y_test.csv", header=None))

In [272]: class GP:
    def __init__(self, alphaS, b):
        self.alphaS = alphaS
        self.b = b
    def fit(self, X, y):
        self.X = X
        self.y = y
        self.Kn = np.array([[self.Kern(xi, xj) for xj in X] for xi in X])
        #print(self.Kn)
        self.pred_help = np.linalg.inv(self.alphaS*np.identity(self.Kn.shape[0])+self.Kn)
        #print(self.pred_help)

    def predict(self, X):
        KX = self.KX_calc(X)
        KDnX = self.KDnX_calc(X)
        helper_var = np.dot(KDnX, self.pred_help)
        pred_mean = np.dot(helper_var, self.y).reshape((X.shape[0],))
        #print(pred_mean)
        pred_var = np.diag(self.alphaS + KX - np.dot(helper_var, KDnX.T))
        #print(pred_mean.shape)
        #print(pred_var.shape)
        return np.array([pred_mean, pred_var]).T # every row is [mean, var]

    def KDnX_calc(self, X):
        return np.array([[self.Kern(xi, xj) for xj in self.X] for xi in X])
```

```

def KX_calc(self, X):
    return np.array([self.Kern(x,x) for x in X])

def Kern(self, xi, xj):
    return np.exp(np.sum(np.square(xi-xj))/(-1*self.b))

def change_pars(self, alphaS, b):
    self.alphaS = alphaS
    self.b = b

def RMSE(y_test, y_pred, leng):
    return np.sqrt(np.sum((y_test-y_pred)**2)/leng)

```

```

In [282]: def part1():
    bees = [5,7,9,11,13,15]
    alphas = [.1,.2,.3,.4,.5,.6,.7,.8,.9,1]
    gaussian_p = GP(0,0)

    RMSE_arr = []
    for b in bees:
        tmp = []
        for alpha in alphas:
            print("fitting", b, alpha)
            gaussian_p.change_pars(alpha, b)
            gaussian_p.fit(X_train, y_train)
            preds = gaussian_p.predict(X_test)
            means = preds[:,0]
            tmp.append(RMSE(y_test.reshape(y_test.shape[0],), means, y_test.shape[0]))
        RMSE_arr.append(tmp)
    return RMSE_arr

```

```

RMSE_arr = part1()

```

```

fitting 5 0.1
fitting 5 0.2
fitting 5 0.3
fitting 5 0.4
fitting 5 0.5
fitting 5 0.6
fitting 5 0.7
fitting 5 0.8
fitting 5 0.9
fitting 5 1
fitting 7 0.1
fitting 7 0.2

```

fitting 7 0.3
fitting 7 0.4
fitting 7 0.5
fitting 7 0.6
fitting 7 0.7
fitting 7 0.8
fitting 7 0.9
fitting 7 1
fitting 9 0.1
fitting 9 0.2
fitting 9 0.3
fitting 9 0.4
fitting 9 0.5
fitting 9 0.6
fitting 9 0.7
fitting 9 0.8
fitting 9 0.9
fitting 9 1
fitting 11 0.1
fitting 11 0.2
fitting 11 0.3
fitting 11 0.4
fitting 11 0.5
fitting 11 0.6
fitting 11 0.7
fitting 11 0.8
fitting 11 0.9
fitting 11 1
fitting 13 0.1
fitting 13 0.2
fitting 13 0.3
fitting 13 0.4
fitting 13 0.5
fitting 13 0.6
fitting 13 0.7
fitting 13 0.8
fitting 13 0.9
fitting 13 1
fitting 15 0.1
fitting 15 0.2
fitting 15 0.3
fitting 15 0.4
fitting 15 0.5
fitting 15 0.6
fitting 15 0.7
fitting 15 0.8
fitting 15 0.9
fitting 15 1

```
In [281]: pdd = pd.DataFrame(RMSE_arr)
pdd.columns = [.1,.2,.3,.4,.5,.6,.7,.8,.9,1]
pdd.reset_index()
pdd.index = [5,7,9,11,13,15]
pdd
```

```
Out[281]:
```

	0.1	0.2	0.3	0.4	0.5	0.6	0.7 \
5	1.966276	1.933135	1.923420	1.922198	1.924769	1.929213	1.934634
7	1.920163	1.904877	1.908080	1.915902	1.924804	1.933701	1.942254
9	1.897649	1.902519	1.917648	1.932514	1.945699	1.957235	1.967403
11	1.890507	1.914981	1.938849	1.957936	1.973216	1.985764	1.996375
13	1.895849	1.935586	1.964597	1.985502	2.001314	2.013878	2.024310
15	1.909603	1.959549	1.990804	2.011915	2.027370	2.039465	2.049463

	0.8	0.9	1.0
5	1.940583	1.946820	1.953213
7	1.950380	1.958093	1.965438
9	1.976492	1.984741	1.992341
11	2.005603	2.013835	2.021345
13	2.033307	2.041317	2.048642
15	2.058105	2.065845	2.072976

1.1.1 C

The best value of RMSE is 1.890507 which is when $b = 11$ and $\alpha_{\text{squared}} = .1$

Here, we change the variance and kernel width. We notice that a low variance ($\alpha_{\text{sqr}} = 1$) works best in general the true underlying Gaussian distribution of our model has a low variance. We also notice that as we change the kernel width (which the larger it is the more we want close values to have a high kernel value) the RMSE converges towards its optimal value ($b=11$) from both sides.

Now, compared to HW1 in which we used least squares ($\lambda=0$) and got an RMSE value of around 2.6 and also ridge regression with $p=2$ and $\lambda=23$ we got an RMSE around 2.2

Clearly, gaussian processes with the kernel we used outperform those models however calculating the inverse of variance+Kn is computationally heavier than LS and RR.

1.1.2 D

```
In [295]: gaussian_p = GP(2,5)
gaussian_p.fit(X_train[:,3], y_train)


train_preds = gaussian_p.predict(X_train[:,3])
train_means = train_preds[:,0]
train_acc = RMSE(y_train.reshape(y_train.shape[0],), train_means, y_train.shape[0])

test_preds = gaussian_p.predict(X_test[:,3])
test_means = test_preds[:,0]
test_acc = RMSE(y_test.reshape(y_test.shape[0],), test_means, y_test.shape[0])
```

```
print("train RMSE: ",train_acc)
print("test RMSE: ",test_acc)
```

```
train RMSE:  4.253743788
test RMSE:  3.40888298673
```

```
In [317]: plt.figure(figsize=(15, 10))
plt.scatter(X_train[:,3], y_train)
order = np.argsort(X_train[:,3])
xs = np.array(X_train[:,3])[order]
ys = np.array(train_means)[order]
plt.plot(xs, ys, color="red")
plt.show()
```



1.2 PART 2

```
In [318]: X_train_B = np.array(pd.read_csv("./boosting/X_train.csv", header=None))
X_test_B = np.array(pd.read_csv("./boosting/X_test.csv", header=None))
y_train_B = np.array(pd.read_csv("./boosting/y_train.csv", header=None))
y_test_B = np.array(pd.read_csv("./boosting/y_test.csv", header=None))
```

```
In [571]: class LS_classifier:
def fit(self, X, y):
```

```

        self.X = X
        self.y = y
        self.w = np.linalg.inv(np.dot(X.T, X)).dot(X.T).dot(self.y)

    def predict(self, X, use_alpha=False):
        predictions = np.sign(np.dot(X, self.w))
        if use_alpha:
            return self.alpha*predictions
        return predictions

    def set_alpha(self, alpha):
        self.alpha = alpha

    def get_alpha(self):
        return self.alpha

    def flipw(self):
        self.w = (-1)*self.w

class Booster:
    def __init__(self, X, y, its, Xtest, ytest):
        self.X = np.hstack((X, np.ones((X.shape[0], 1))))
        self.y = y.reshape(y.shape[0],)
        self.Xtest = np.hstack((Xtest, np.ones((Xtest.shape[0], 1))))
        self.ytest = ytest.reshape(ytest.shape[0],)
        self.Xb = self.X
        self.yb = self.y
        self.its = its
        self.weights = np.array([1/X.shape[0] for i in range(X.shape[0])])
        self.classifiers = []
        self.error_arr = []
        self.test_er = []
        self.train_er = []
        self.train_uppab = []
        self.boots = [0 for i in range(self.X.shape[0])]
    def boost(self):
        for t in range(self.its):
            if (t+1)%300==0:
                print("iteration: ",t+1)
            #         if t<10:
            #             print(self.weights)
            self.bootstrap()
            tmp_classif = LS_classifier()
            tmp_classif.fit(self.Xb, self.yb)
            preds = tmp_classif.predict(self.X)
            #check
            #print(preds.shape==self.y.shape)
            #error = np.dot(self.weights,np.array(preds!=self.y, dtype=int).T)

```

```

#EPSILON
error = np.sum(self.weights[(preds!=self.y)])
if error>0.5:
    tmp_classif.flipw()
    preds = tmp_classif.predict(self.X)
    #error = np.dot(self.weights,np.array(preds!=self.y, dtype=int).T)
    error = np.sum(self.weights[(preds!=self.y)])
self.error_arr.append(error)

#print("error ",error)
alpha = np.log((1-error)/error)*0.5
#print("alpha ",alpha)
tmp_classif.set_alpha(alpha)
#self.weights *= np.exp(-alpha*np.dot(self.y, preds))
self.weights = normalize(np.multiply(np.exp(-alpha*np.multiply(self.y,preds))
#print("weights ", self.weights)
#print(np.exp(-alpha*np.dot(self.y, preds)))
#self.norm_w()
self.classifiers.append(tmp_classif)

# TRAIN ERROR
train_preds = self.predict(self.X)
train_error = errorcalc(self.y, train_preds, train_preds.shape[0])
self.train_er.append(train_error)

# TEST ERROR
test_preds = self.predict(self.Xtest)
test_error = errorcalc(self.ytest, test_preds, test_preds.shape[0])
self.test_er.append(test_error)

# TRAIN UPPERBOUND
ZT = np.exp(-alpha)*(1-error) + np.exp(alpha)*error
if len(self.train_uppab) == 0:
    self.train_uppab.append(ZT)
else:
    self.train_uppab.append(ZT*self.train_uppab[-1])

def predict(self, X):
    res = []
    for model in self.classifiers:
        res.append(model.predict(X, True))
    res = np.array(res)
    return np.sign(np.sum(res, axis=0))

```

```

def bootstrap(self):
    indx = np.random.choice(np.arange(self.X.shape[0]), size=self.X.shape[0], replace=True)
    #self.Xb = np.take(self.X, indices=indx, axis=0)
    #self.yb = np.take(self.y, indices=indx)
    for i in indx:
        self.boots[i] += 1
    self.Xb = self.X[indx]
    self.yb = self.y[indx]

def get_error(self):
    return self.error_arr, self.train_er, self.test_er, self.train_ensemble, self.test_ensemble

def get_alphas(self):
    res = []
    for model in self.classifiers:
        res.append(model.get_alpha())
    return res

def errorcalc(x, y, l):
    return np.sum(np.array(x==y, dtype=int))/l

def normalize(vec):
    return vec/np.sum(vec)

```

```

In [578]: boosty = Booster(X_train_B, y_train_B, 1500, X_test_B, y_test_B.reshape((y_test_B.shape[0], 1)))
          boosty.boost()

```

```

iteration: 300
iteration: 600
iteration: 900
iteration: 1200
iteration: 1500

```

```


In [579]: errors = boosty.get_error()

```


```

plt.figure(figsize=(15, 10))
plt.plot(range(1,1501), 1-np.array(errors[1]), label= "train error")
plt.plot(range(1,1501), 1-np.array(errors[2]), label= "test error")
plt.legend()
plt.show()


```


```
In [574]: plt.figure(figsize=(15, 10))  
          plt.plot(range(0,1500), np.array(errors[3]), label= "upperbound")  
          plt.legend()  
          plt.show()
```



```
In [575]: plt.figure(figsize=(20, 10))
          markerline, stemlines, baseline = plt.stem(range(np.array(errors[4]).shape[0]), np.array(errors[4]))
          plt.setp(markerline, visible=False)
          plt.setp(stemlines, linewidth=1)
          plt.setp(baseline, visible=False)
          plt.legend()
          plt.show()
```



```
In [576]: plt.figure(figsize=(15, 10))
plt.plot(range(0,1500), np.array(errors[0]), label= "epsilon")
plt.legend()
plt.show()
```



```
In [577]: plt.figure(figsize=(15, 10))
plt.plot(range(0,1500), boosty.get_alphas(), label= "alpha")
plt.legend()
plt.show()
```

